Building the Trust Layer for Autonomous AI

AI agents are becoming autonomous.

Every autonomous agent needs governance.

Seed Round | $1.5M | July 2026

AYMAN SEIF
Founder & CEO

CURRENT STATUS

Company

ELMAHROSA International

Egypt Ø<ÝêØ<Ýì
MENA-focused AI Security Company

Registered Member, Anthropic Claude Partner Network

Programs

TEOS Sentinel Shield — AI execution firewall, pre-execution governance

TEOS AI Engine — Multi-model orchestration (Groq, Claude, OpenAI)

TEOS Sovereign Security Stack — Policy engine, MCP, Railway production infra

Technology

121+ Security Rules | Pre-Execution AI Validation

Deterministic Risk Engine | MCP Integration

Railway Production Infrastructure | Redis | PostgreSQL

THE PROBLEM

Autonomous Systems Without Governance

Organizations are deploying:

- AI Agents — autonomous executors making production decisions
- MCP Servers — model-context-protocol servers with system access
- Autonomous Workflows — multi-step AI pipelines without human oversight
- Code Generation Systems — AI code running in production environments

TEOS focuses on: Prevent Before Execution

- Secret Exposure
- Supply Chain Risk
- Data Exfiltration
- Dangerous Commands
- Prompt Injection
- Policy Violations

MARKET OPPORTUNITY

Verified Market Signals

These are publicly reported trends from leading analyst firms. Not projections — signals.

GARTNER

Enterprise AI spending continues to grow globally. AI Governance is emerging as a mandatory procurement category.

IDC

AI Governance spending is projected to grow as organizations deploying AI require auditability and runtime controls.

MCKINSEY

Governments worldwide are introducing AI regulations (EU AI Act, US Executive Order) requiring governance infrastructure.

DELOITTE

Enterprises need auditability and execution controls for AI agents. Without governance, autonomous systems introduce unacceptable risk.

Sources: Gartner "AI Governance Market Forecast" | IDC "Worldwide AI Spending Guide" | McKinsey "State of AI" | Deloitte "AI Risk Management"

OUR SOLUTION

Governance Before Execution

TEOS Sentinel Shield is a runtime policy engine for AI agents.

It prevents dangerous actions before they execute.

01

Detect

121+ rules analyze agent behavior, code, and commands in real time before execution

02

Block

Policy engine prevents dangerous actions — binary ALLOW/BLOCK in under 100ms

03

Report

Audit trails, compliance reports, and governance logs for CISOs

Key Differentiator

Most security tools audit after execution. We block before damage occurs.

PRODUCT

Enterprise Power, Three Surfaces

Telegram Bot

Scan code, commands, or tokens by pasting. Zero learning curve, enterprise power.

Enforcement API

RESTful API for enterprise integration. Real-time policy callbacks.

PDF Reports

Due diligence reports with risk gauges, analysis, and audit trails.

MCP Integration

Native Model-Context-Protocol support. Govern at the protocol level.

Policy Engine

Configurable policies — block, flag, or allow based on severity.

Audit Trail

Immutable governance logs. Every scan, every decision, timestamped.

TECHNOLOGY

Built for Scale, Engineered for Trust

121+ Governance Rules	Deterministic engine — no LLM latency, no hallucination, binary decisions.
Pre-Execution Validation	Actions classified and blocked before execution. Preventive, not reactive.
Multi-Model AI	Groq, Claude, OpenAI — automatic failover, cost-optimized routing.
MCP Integration	Native protocol-level governance for AI agents.

INFRASTRUCTURE

Node.js 24 | PDFKit | Redis | PostgreSQL | Railway Cloud | Telegram Bot API | Solana RPC | Jupiter/Birdeye | GitHub API

COMPETITIVE POSITIONING

Why TEOS Wins

	Traditional Security	AI Code Assistants	TEOS Sentinel
Approach	Reactive auditing	Generates code	
Focus	Human operators	Developers	
Timing	After incidents	At development	
Enforcement	Manual review	Suggestions only	

TEOS is the only platform that governs AI agent execution before it happens.

Competitors audit after the fact. We block before damage occurs.

DEFENSIBLE ADVANTAGES

Why We Win

Deterministic Policy Engine

121 rules in <100ms. No LLM latency, no hallucination.

Pre-Execution Enforcement

Blocks BEFORE damage. Preventive, not post-incident.

Multi-Model AI

Groq + Claude + OpenAI. Auto failover. No single vendor.

Enterprise Audit Trail

Immutable logs, compliance reports, CISO-ready.

Multi-Channel

Telegram + API + MCP. Same engine, three surfaces.

Anthropic Partner

Registered member, Claude integration, co-marketing.

The Moat Deepens With Scale

More scans !' better rules | More data !' better AI | More integrations !' switching costs

BUSINESS MODEL

SaaS + Services

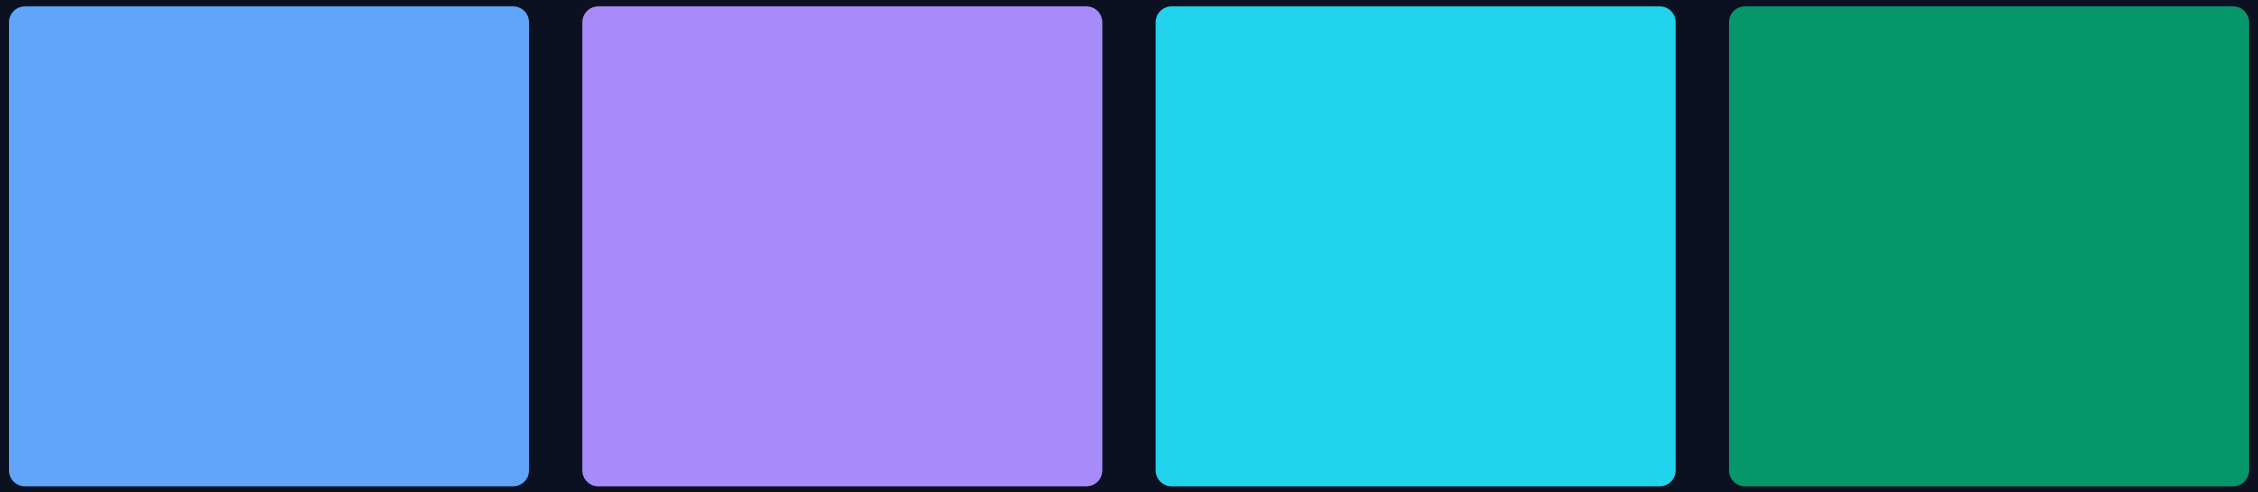

Additional Revenue

Deployment Services

Compliance Packages

Air-Gapped Installations

Security Assessments

Training

CURRENT TRACTION

What We Have Built

Only what is true. No fabricated numbers.

Product

- Railway Deployment — production infrastructure active
- MCP Infrastructure — protocol integration live
- Telegram Operations Bot — scanning and governance bot active
- Security Rule Engine — 121+ rules, deterministic enforcement
- Production Checkout System — billing and payments integrated

Community

- Active Development Community — ongoing contributions
- MENA Focus — regional presence and engagement
- Global Early Access Program — beta users onboarded

Partnerships

- Anthropic Claude Partner Network — registered member

THE ASK

Seed Round: $1.5M

Building the governance layer for autonomous AI.

Use of Funds

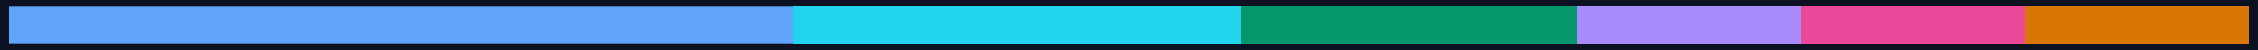

- **Engineering $525K (35%)**
- **Enterprise Sales $300K (20%)**
- **Security & Compliance $225K (15%)**
- **Cloud Infrastructure $150K (10%)**
- **Marketing $150K (10%)**
- **Working Capital $150K (10%)**

12-Month Targets

These are goals, not actuals.

100 Paying Teams

10 Enterprise Pilots

1,000 Developer Accounts

3 Government Pilots

TEAM

Solo Founder, Full-Stack Builder

Ayman Seif

Founder & CEO | Sole Engineer

Serial entrepreneur. 10K+ users across MENA.
Full-stack: Node.js, Python, React, Solidity.
Built v4.0 GA solo (121 rules, PDF reports,
billing, bot, API, dashboard).
Deep expertise in blockchain, AI/ML,
distributed systems.

Advisory (Recruiting)

- CTO Advisor — DeFi/security, enterprise AI experience
- Board Advisor — Former CISO at fintech company
- University partnerships for AI model research & talent

RISKS & MITIGATION

We Know the Risks

RISK: Regulatory uncertainty in AI governance

MITIGATION: Governance-first architecture. Multi-vertical positioning avoids single-regulation dependency.

RISK: Competition from well-funded vendors

MITIGATION: First-mover in pre-execution governance. Deep technical moat. Deterministic engine is hard to replicate.

RISK: Single founder risk

MITIGATION: Aggressive hiring post-seed. Advisory board underway. Product already in production (v4.0 GA).

RISK: Market adoption slower than projected

MITIGATION: Lean burn rate extends runway. Multiple go-to-market channels (Telegram, API, MCP).

VISION

The Future We Are Building

Every enterprise will deploy AI agents.

Every AI agent will require governance.

TEOS Sentinel Shield is building

the trust infrastructure for autonomous AI.

From code security to enterprise AI governance.

The same engine. The same trust layer.

Every Agent Will Need Governance.

We are building the trust layer for autonomous AI.

Ayman Seif | Founder & CEO

ayman@teosegypt.com

Raise: $1.5M Seed | Sent from Cairo, Shipping Globally

https://sentinel.teosegypt.com | https://github.com/Elmahrosa